

bounding box

May 19, 2019

0.0.1 Read the cropping dataset

```
In [1]: import keras
```

```
with open('cropping.txt', 'rt') as f: data = f.read().split('\n')[:-1]
data = [line.split(',') for line in data]
data = [(p, [(int(coord[i]), int(coord[i+1])) for i in range(0, len(coord), 2)]) for p, *coo
data[0]
```

```
/anaconda3/lib/python3.6/site-packages/h5py/_init_.py:36: FutureWarning: Conversion of the s
from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

```
Out[1]: ('88532e70.jpg',
        [(195, 293),
         (269, 115),
         (868, 158),
         (888, 170),
         (641, 496),
         (512, 546),
         (321, 524)])
```

This dataset contains the location of points on the edge of the fluke for 1200 pictures randomly selected from the Humpback Whale Identification Challenge training set. Points are selected to capture the leftmost and rightmost points, as well as the highest and lowest points. Additional points are added to help determine the fluke bounding box following an affine transformation on the image.

The intent of this dataset is to build a model for locating the whale fluke inside the image. In the context of the Humpback Whale Identification Challenge, such a model can then be used to crop images around the region of interest.

```
In [2]: from PIL import Image as pil_image
        from PIL.ImageDraw import Draw
        from os.path import isfile

        def expand_path(p):
            if isfile('train/' + p): return 'train/' + p
```

```

    if isfile('test/' + p): return 'test/' + p
    return p

def read_raw_image(p):
    return pil_image.open(expand_path(p))

def draw_dot(draw, x, y):
    draw.ellipse(((x-5,y-5),(x+5,y+5)), fill='red', outline='red')

def draw_dots(draw, coordinates):
    for x,y in coordinates: draw_dot(draw, x, y)

def bounding_rectangle(list):
    x0, y0 = list[0]
    x1, y1 = x0, y0
    for x,y in list[1:]:
        x0 = min(x0, x)
        y0 = min(y0, y)
        x1 = max(x1, x)
        y1 = max(y1, y)
    return x0,y0,x1,y1

filename,coordinates = data[0]
box = bounding_rectangle(coordinates)
img = read_raw_image(filename)
draw = Draw(img)
draw_dots(draw, coordinates)
draw.rectangle(box, outline='red')
img

```

Out[2]:



```
In [3]: # Define useful constants
img_shape = (128,128,1)
anisotropy = 2.15
```

```
In [4]: import random
import numpy as np
from scipy.ndimage import affine_transform
from keras.preprocessing.image import img_to_array
```

```
# Read an image as black&white numpy array
```

```
def read_array(p):
    img = read_raw_image(p).convert('L')
    return img_to_array(img)
```

```
def build_transform(rotation, shear, height_zoom, width_zoom, height_shift, width_shift):
    rotation = np.deg2rad(rotation)
    shear = np.deg2rad(shear)
    rotation_matrix = np.array([[np.cos(rotation), np.sin(rotation), 0], [-np.sin(rotation), np.cos(rotation), 0], [0, 0, 1]])
    shift_matrix = np.array([[1, 0, height_shift], [0, 1, width_shift], [0, 0, 1]])
    shear_matrix = np.array([[1, np.sin(shear), 0], [0, np.cos(shear), 0], [0, 0, 1]])
    zoom_matrix = np.array([[1.0/height_zoom, 0, 0], [0, 1.0/width_zoom, 0], [0, 0, 1]])
    shift_matrix = np.array([[1, 0, -height_shift], [0, 1, -width_shift], [0, 0, 1]])
    return np.dot(np.dot(rotation_matrix, shear_matrix), np.dot(zoom_matrix, shift_matrix))
```

```

# Compute the coordinate transformation required to center the pictures, padding as required
def center_transform(affine, input_shape):
    hi, wi = float(input_shape[0]), float(input_shape[1])
    ho, wo = float(img_shape[0]), float(img_shape[1])
    top, left, bottom, right = 0, 0, hi, wi
    if wi/hi/anisotropy < wo/ho: # input image too narrow, extend width
        w = hi*wo/ho*anisotropy
        left = (wi-w)/2
        right = left + w
    else: # input image too wide, extend height
        h = wi*ho/wo/anisotropy
        top = (hi-h)/2
        bottom = top + h
    center_matrix = np.array([[1, 0, -ho/2], [0, 1, -wo/2], [0, 0, 1]])
    scale_matrix = np.array([(bottom - top)/ho, 0, 0], [0, (right - left)/wo, 0],
    decenter_matrix = np.array([[1, 0, hi/2], [0, 1, wi/2], [0, 0, 1]])
    return np.dot(np.dot(decenter_matrix, scale_matrix), np.dot(affine, center_matrix))

# Apply an affine transformation to an image represented as a numpy array.
def transform_img(x, affine):
    matrix = affine[:2,:2]
    offset = affine[:2,2]
    x = np.moveaxis(x, -1, 0)
    channels = [affine_transform(channel, matrix, offset, output_shape=img_shape[:-1],
    mode='constant', cval=np.average(channel)) for channel in x]
    return np.moveaxis(np.stack(channels, axis=0), 0, -1)

# Read an image for validation, i.e. without data augmentation.
def read_for_validation(p):
    x = read_array(p)
    t = np.array([[1, 0, 0], [0, 1, 0], [0, 0, 1]])
    t = center_transform(t, x.shape)
    x = transform_img(x, t)
    x -= np.mean(x, keepdims=True)
    x /= np.std(x, keepdims=True) + K.epsilon()
    return x,t

# Read an image for training, i.e. including a random affine transformation
def read_for_training(p):
    x = read_array(p)
    t = build_transform(
        random.uniform(-5, 5),
        random.uniform(-5, 5),
        random.uniform(0.9, 1.0),
        random.uniform(0.9, 1.0),
        random.uniform(-0.05*img_shape[0], 0.05*img_shape[0]),
        random.uniform(-0.05*img_shape[1], 0.05*img_shape[1]))
    t = center_transform(t, x.shape)

```

```

x = transform_img(x, t)
x -= np.mean(x, keepdims=True)
x /= np.std(x, keepdims=True) + K.epsilon()
return x,t

# Transform corrdinates according to the provided affine transformation
def coord_transform(list, trans):
    result = []
    for x,y in list:
        y,x,_ = trans.dot([y,x,1]).astype(np.int)
        result.append((x,y))
    return result

```

In [5]: `from sklearn.model_selection import train_test_split`

```

train, val = train_test_split(data, test_size=200, random_state=1)
train += train
train += train
train += train
train += train
len(train),len(val)

```

Out [5]: (16000, 200)

In [7]: `import matplotlib.pyplot as plt`

```

from tqdm import tqdm, tqdm_notebook
from keras import backend as K
from keras.preprocessing.image import array_to_img
from numpy.linalg import inv as mat_inv

def show_whale(imgs, per_row=5):
    n = len(imgs)
    rows = (n + per_row - 1)//per_row
    cols = min(per_row, n)
    fig, axes = plt.subplots(rows,cols, figsize=(24//per_row*cols,24//per_row*rows))
    for ax in axes.flatten(): ax.axis('off')
    for i,(img,ax) in enumerate(zip(imgs, axes.flatten())): ax.imshow(img.convert('RGB'))

val_a = np.zeros((len(val),)+img_shape,dtype=K.floatx()) # Preprocess validation image
val_b = np.zeros((len(val),4),dtype=K.floatx()) # Preprocess bounding boxes
for i,(p,coords) in enumerate(tqdm_notebook(val)):
    img,trans = read_for_validation(p)
    coords = coord_transform(coords, mat_inv(trans))
    x0,y0,x1,y1 = bounding_rectangle(coords)
    val_a[i,:,:,:] = img
    val_b[i,0] = x0
    val_b[i,1] = y0
    val_b[i,2] = x1

```


```

        val_b[i,3]      = y1

    idx = 1
    img = array_to_img(val_a[idx])
    img = img.convert('RGB')
    draw = Draw(img)
    draw.rectangle(val_b[idx], outline='red')
    show_whale([read_raw_image(val[idx][0]), img], per_row=2)

HBox(children=(IntProgress(value=0, max=200), HTML(value='')))

```



```
In [8]: from keras.utils import Sequence
```

```

class TrainingData(Sequence):
    def __init__(self, batch_size=32):
        super(TrainingData, self).__init__()
        self.batch_size = batch_size
    def __getitem__(self, index):
        start = self.batch_size*index;
        end = min(len(train), start + self.batch_size)
        size = end - start
        a = np.zeros((size,) + img_shape, dtype=K.floatx())
        b = np.zeros((size,4), dtype=K.floatx())
        for i,(p,coords) in enumerate(train[start:end]):
            img,trans = read_for_training(p)
            coords = coord_transform(coords, mat_inv(trans))


```

```

        x0,y0,x1,y1 = bounding_rectangle(coords)
        a[i,:,:,:] = img
        b[i,0]      = x0
        b[i,1]      = y0
        b[i,2]      = x1
        b[i,3]      = y1
    return a,b
def __len__(self):
    return (len(train) + self.batch_size - 1)//self.batch_size

random.seed(1)
a, b = TrainingData(batch_size=5)[1]
img = array_to_img(a[0])
img = img.convert('RGB')
draw = Draw(img)
draw.rectangle(b[0], outline='red')
show_whale([read_raw_image(train[0][0]), img], per_row=2)

```



```

In [9]: from keras.engine.topology import Input
        from keras.layers import BatchNormalization, Concatenate, Conv2D, Dense, Dropout, Flatten
        from keras.models import Model

        def build_model(with_dropout=True):
            kwargs = {'activation':'relu', 'padding':'same'}
            conv_drop = 0.2
            dense_drop = 0.5
            inp = Input(shape=img_shape)

            x = inp

```

```

x = Conv2D(64, (9, 9), **kwargs)(x)
x = Conv2D(64, (3, 3), **kwargs)(x)
x = BatchNormalization()(x)
if with_dropout: x = Dropout(conv_drop, noise_shape=(None, 1, 1, int(x.shape[-1])))

x = Conv2D(64, (2, 2), **kwargs, strides=2)(x)
x = Conv2D(64, (3, 3), **kwargs)(x)
x = Conv2D(64, (3, 3), **kwargs)(x)
x = BatchNormalization()(x)
if with_dropout: x = Dropout(conv_drop, noise_shape=(None, 1, 1, int(x.shape[-1])))

x = Conv2D(64, (2, 2), **kwargs, strides=2)(x)
x = Conv2D(64, (3, 3), **kwargs)(x)
x = Conv2D(64, (3, 3), **kwargs)(x)
x = BatchNormalization()(x)
if with_dropout: x = Dropout(conv_drop, noise_shape=(None, 1, 1, int(x.shape[-1])))

x = Conv2D(64, (2, 2), **kwargs, strides=2)(x)
x = Conv2D(64, (3, 3), **kwargs)(x)
x = Conv2D(64, (3, 3), **kwargs)(x)
x = BatchNormalization()(x)
if with_dropout: x = Dropout(conv_drop, noise_shape=(None, 1, 1, int(x.shape[-1])))

x = Conv2D(64, (2, 2), **kwargs, strides=2)(x)
x = Conv2D(64, (3, 3), **kwargs)(x)
x = Conv2D(64, (3, 3), **kwargs)(x)
x = BatchNormalization()(x)
if with_dropout: x = Dropout(conv_drop, noise_shape=(None, 1, 1, int(x.shape[-1])))

h = MaxPooling2D(pool_size=(1, int(x.shape[2])))(x)
h = Flatten()(h)
if with_dropout: h = Dropout(dense_drop)(h)
h = Dense(16, activation='relu')(h)

v = MaxPooling2D(pool_size=(int(x.shape[1]), 1))(x)
v = Flatten()(v)
if with_dropout: v = Dropout(dense_drop)(v)
v = Dense(16, activation='relu')(v)

x = Concatenate()([h,v])
if with_dropout: x = Dropout(0.5)(x)
x = Dense(4, activation='linear')(x)

```



```

return Model(inp,x)

model = build_model(with_dropout=True)
model.summary()

```

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	(None, 128, 128, 1)	0	
conv2d_1 (Conv2D)	(None, 128, 128, 64)	5248	input_1[0][0]
conv2d_2 (Conv2D)	(None, 128, 128, 64)	36928	conv2d_1[0][0]
batch_normalization_1 (BatchNor	(None, 128, 128, 64)	256	conv2d_2[0][0]
dropout_1 (Dropout)	(None, 128, 128, 64)	0	batch_normalization_1[0][0]
conv2d_3 (Conv2D)	(None, 64, 64, 64)	16448	dropout_1[0][0]
conv2d_4 (Conv2D)	(None, 64, 64, 64)	36928	conv2d_3[0][0]
conv2d_5 (Conv2D)	(None, 64, 64, 64)	36928	conv2d_4[0][0]
batch_normalization_2 (BatchNor	(None, 64, 64, 64)	256	conv2d_5[0][0]
dropout_2 (Dropout)	(None, 64, 64, 64)	0	batch_normalization_2[0][0]
conv2d_6 (Conv2D)	(None, 32, 32, 64)	16448	dropout_2[0][0]
conv2d_7 (Conv2D)	(None, 32, 32, 64)	36928	conv2d_6[0][0]
conv2d_8 (Conv2D)	(None, 32, 32, 64)	36928	conv2d_7[0][0]
batch_normalization_3 (BatchNor	(None, 32, 32, 64)	256	conv2d_8[0][0]
dropout_3 (Dropout)	(None, 32, 32, 64)	0	batch_normalization_3[0][0]
conv2d_9 (Conv2D)	(None, 16, 16, 64)	16448	dropout_3[0][0]
conv2d_10 (Conv2D)	(None, 16, 16, 64)	36928	conv2d_9[0][0]
conv2d_11 (Conv2D)	(None, 16, 16, 64)	36928	conv2d_10[0][0]
batch_normalization_4 (BatchNor	(None, 16, 16, 64)	256	conv2d_11[0][0]
dropout_4 (Dropout)	(None, 16, 16, 64)	0	batch_normalization_4[0][0]

conv2d_12 (Conv2D)	(None, 8, 8, 64)	16448	dropout_4[0][0]
conv2d_13 (Conv2D)	(None, 8, 8, 64)	36928	conv2d_12[0][0]
conv2d_14 (Conv2D)	(None, 8, 8, 64)	36928	conv2d_13[0][0]
batch_normalization_5 (BatchNor	(None, 8, 8, 64)	256	conv2d_14[0][0]
dropout_5 (Dropout)	(None, 8, 8, 64)	0	batch_normalization_5[0][0]
conv2d_15 (Conv2D)	(None, 4, 4, 64)	16448	dropout_5[0][0]
conv2d_16 (Conv2D)	(None, 4, 4, 64)	36928	conv2d_15[0][0]
conv2d_17 (Conv2D)	(None, 4, 4, 64)	36928	conv2d_16[0][0]
batch_normalization_6 (BatchNor	(None, 4, 4, 64)	256	conv2d_17[0][0]
dropout_6 (Dropout)	(None, 4, 4, 64)	0	batch_normalization_6[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 4, 1, 64)	0	dropout_6[0][0]
max_pooling2d_2 (MaxPooling2D)	(None, 1, 4, 64)	0	dropout_6[0][0]
flatten_1 (Flatten)	(None, 256)	0	max_pooling2d_1[0][0]
flatten_2 (Flatten)	(None, 256)	0	max_pooling2d_2[0][0]
dropout_7 (Dropout)	(None, 256)	0	flatten_1[0][0]
dropout_8 (Dropout)	(None, 256)	0	flatten_2[0][0]
dense_1 (Dense)	(None, 16)	4112	dropout_7[0][0]
dense_2 (Dense)	(None, 16)	4112	dropout_8[0][0]
concatenate_1 (Concatenate)	(None, 32)	0	dense_1[0][0] dense_2[0][0]
dropout_9 (Dropout)	(None, 32)	0	concatenate_1[0][0]
dense_3 (Dense)	(None, 4)	132	dropout_9[0][0]
Total params: 503,588			
Trainable params: 502,820			
Non-trainable params: 768			

```

In [10]: from keras.callbacks import EarlyStopping, ModelCheckpoint, ReduceLROnPlateau
         from keras.optimizers import Adam

         for num in range(1, 4):
             model_name = 'cropping-%01d.h5' % num
             print(model_name)
             model.compile(Adam(lr=0.032), loss='mean_squared_error')
             model.fit_generator(
                 TrainingData(), epochs=50, max_queue_size=12, workers=4, verbose=1,
                 validation_data=(val_a, val_b),
                 callbacks=[
                     EarlyStopping(monitor='val_loss', patience=9, min_delta=0.1, verbose=1),
                     ReduceLROnPlateau(monitor='val_loss', patience=3, min_delta=0.1, factor=0.5),
                     ModelCheckpoint(model_name, save_best_only=True, save_weights_only=True),
                 ])
             model.load_weights(model_name)
             model.evaluate(val_a, val_b, verbose=0)

cropping-1.h5
Epoch 1/50
500/500 [=====] - 2507s 5s/step - loss: 557.0855 - val_loss: 165.9808
Epoch 2/50
500/500 [=====] - 2540s 5s/step - loss: 346.6362 - val_loss: 120.6931
Epoch 3/50
500/500 [=====] - 2574s 5s/step - loss: 278.8406 - val_loss: 52731.49
Epoch 4/50
500/500 [=====] - 2575s 5s/step - loss: 240.7661 - val_loss: 131.0033
Epoch 5/50
500/500 [=====] - 2531s 5s/step - loss: 208.5899 - val_loss: 12988.34

Epoch 00005: ReduceLROnPlateau reducing learning rate to 0.00800000037997961.
Epoch 6/50
500/500 [=====] - 2482s 5s/step - loss: 179.2290 - val_loss: 68.9442
Epoch 7/50
500/500 [=====] - 2388s 5s/step - loss: 170.2741 - val_loss: 70.3035
Epoch 8/50
500/500 [=====] - 2289s 5s/step - loss: 162.0193 - val_loss: 68.0217
Epoch 9/50
500/500 [=====] - 2284s 5s/step - loss: 152.7652 - val_loss: 56.9733
Epoch 10/50
500/500 [=====] - 2333s 5s/step - loss: 144.5138 - val_loss: 55.2711
Epoch 11/50
500/500 [=====] - 2452s 5s/step - loss: 139.8426 - val_loss: 48.1526
Epoch 12/50
500/500 [=====] - 2447s 5s/step - loss: 131.9710 - val_loss: 54.0475
Epoch 13/50
500/500 [=====] - 2380s 5s/step - loss: 122.8758 - val_loss: 47.9916
Epoch 14/50

```

500/500 [=====] - 2376s 5s/step - loss: 116.9095 - val_loss: 50.3715
Epoch 15/50
500/500 [=====] - 2382s 5s/step - loss: 113.2557 - val_loss: 45.1115
Epoch 16/50
500/500 [=====] - 2382s 5s/step - loss: 105.9497 - val_loss: 40.0823
Epoch 17/50
500/500 [=====] - 2381s 5s/step - loss: 99.2378 - val_loss: 50.9979
Epoch 18/50
500/500 [=====] - 2384s 5s/step - loss: 92.9641 - val_loss: 52.6463
Epoch 19/50
500/500 [=====] - 2376s 5s/step - loss: 89.5002 - val_loss: 40.0454

Epoch 00019: ReduceLROnPlateau reducing learning rate to 0.0020000000949949026.

Epoch 20/50
500/500 [=====] - 2384s 5s/step - loss: 83.6473 - val_loss: 36.9560
Epoch 21/50
500/500 [=====] - 2382s 5s/step - loss: 79.4791 - val_loss: 34.4707
Epoch 22/50
500/500 [=====] - 2383s 5s/step - loss: 79.4619 - val_loss: 30.8307
Epoch 23/50
500/500 [=====] - 2381s 5s/step - loss: 77.7097 - val_loss: 34.4418
Epoch 24/50
500/500 [=====] - 2383s 5s/step - loss: 76.6076 - val_loss: 31.3798
Epoch 25/50
500/500 [=====] - 2361s 5s/step - loss: 75.6522 - val_loss: 32.4307

Epoch 00025: ReduceLROnPlateau reducing learning rate to 0.002.

Epoch 26/50
500/500 [=====] - 2385s 5s/step - loss: 73.5478 - val_loss: 31.1684
Epoch 27/50
500/500 [=====] - 2388s 5s/step - loss: 73.4404 - val_loss: 29.9076
Epoch 28/50
500/500 [=====] - 2385s 5s/step - loss: 71.9389 - val_loss: 32.5809
Epoch 29/50
500/500 [=====] - 2390s 5s/step - loss: 71.8188 - val_loss: 29.5498
Epoch 30/50
500/500 [=====] - 2388s 5s/step - loss: 70.3340 - val_loss: 32.6804
Epoch 31/50
500/500 [=====] - 2406s 5s/step - loss: 67.5038 - val_loss: 36.6804
Epoch 32/50
500/500 [=====] - 2411s 5s/step - loss: 68.0174 - val_loss: 31.5727

Epoch 00032: ReduceLROnPlateau reducing learning rate to 0.002.

Epoch 33/50
500/500 [=====] - 2395s 5s/step - loss: 67.3362 - val_loss: 33.4936
Epoch 34/50
500/500 [=====] - 2402s 5s/step - loss: 65.4268 - val_loss: 33.1911
Epoch 35/50

```

500/500 [=====] - 2399s 5s/step - loss: 65.7545 - val_loss: 33.0299

Epoch 00035: ReduceLROnPlateau reducing learning rate to 0.002.
Epoch 36/50
500/500 [=====] - 2434s 5s/step - loss: 64.4783 - val_loss: 32.2387
Epoch 37/50
500/500 [=====] - 2432s 5s/step - loss: 63.4011 - val_loss: 29.4551
Epoch 38/50
500/500 [=====] - 2441s 5s/step - loss: 62.9743 - val_loss: 28.3006
Epoch 39/50
500/500 [=====] - 2456s 5s/step - loss: 61.8167 - val_loss: 27.4117
Epoch 40/50
500/500 [=====] - 2470s 5s/step - loss: 61.1573 - val_loss: 31.0975
Epoch 41/50
500/500 [=====] - 2392s 5s/step - loss: 60.0332 - val_loss: 30.8048
Epoch 42/50
500/500 [=====] - 2388s 5s/step - loss: 59.3688 - val_loss: 28.4462

Epoch 00042: ReduceLROnPlateau reducing learning rate to 0.002.
Epoch 43/50
500/500 [=====] - 2396s 5s/step - loss: 58.6502 - val_loss: 26.9193
Epoch 44/50
500/500 [=====] - 2396s 5s/step - loss: 57.9148 - val_loss: 30.2132
Epoch 45/50
500/500 [=====] - 2381s 5s/step - loss: 57.0164 - val_loss: 31.5943
Epoch 46/50
500/500 [=====] - 2438s 5s/step - loss: 56.2615 - val_loss: 28.4853

Epoch 00046: ReduceLROnPlateau reducing learning rate to 0.002.
Epoch 47/50
500/500 [=====] - 2508s 5s/step - loss: 56.1670 - val_loss: 30.7610
Epoch 48/50
500/500 [=====] - 2469s 5s/step - loss: 55.5824 - val_loss: 30.8394
Epoch 49/50
500/500 [=====] - 2594s 5s/step - loss: 54.3359 - val_loss: 29.4281

Epoch 00049: ReduceLROnPlateau reducing learning rate to 0.002.
Epoch 50/50
500/500 [=====] - 2725s 5s/step - loss: 53.9117 - val_loss: 30.5559
cropping-2.h5
Epoch 1/50
354/500 [=====>...] - ETA: 11:46 - loss: 83.0287

```

KeyboardInterrupt

Traceback (most recent call last)

```

<ipython-input-10-f9fc98992254> in <module>()
    12         EarlyStopping(monitor='val_loss', patience=9, min_delta=0.1, verbose=1)
    13         ReduceLROnPlateau(monitor='val_loss', patience=3, min_delta=0.1, factor=0.5)
--> 14         ModelCheckpoint(model_name, save_best_only=True, save_weights_only=True)
    15     ])
    16     model.load_weights(model_name)

/anaconda3/lib/python3.6/site-packages/keras/legacy/interfaces.py in wrapper(*args, **kwargs)
    89         warnings.warn('Update your `' + object_name + '` call to the `keras.' + object_name +
    90             'Keras 2 API: ' + signature, stacklevel=2)
--> 91         return func(*args, **kwargs)
    92     wrapper._original_function = func
    93     return wrapper

/anaconda3/lib/python3.6/site-packages/keras/engine/training.py in fit_generator(self, x_generator,
1416         use_multiprocessing=use_multiprocessing,
1417         shuffle=shuffle,
-> 1418         initial_epoch=initial_epoch)
1419
1420     @interfaces.legacy_generator_methods_support

/anaconda3/lib/python3.6/site-packages/keras/engine/training_generator.py in fit_generator(self, x_generator,
215         outs = model.train_on_batch(x, y,
216             sample_weight=sample_weight,
--> 217             class_weight=class_weight)
218
219         outs = to_list(outs)

/anaconda3/lib/python3.6/site-packages/keras/engine/training.py in train_on_batch(self, x, y, sample_weight,
1215         ins = x + y + sample_weights
1216         self._make_train_function()
-> 1217         outputs = self.train_function(ins)
1218         return unpack_singleton(outputs)
1219

/anaconda3/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py in __call__(self, inputs)
2713         return self._legacy_call(inputs)
2714
-> 2715         return self._call(inputs)
2716     else:
2717         if py_any(is_tensor(x) for x in inputs):

```

```

/anaconda3/lib/python3.6/site-packages/keras/backend/tensorflow_backend.py in _call(se
2673         fetched = self._callable_fn(*array_vals, run_metadata=self.run_metadata)
2674     else:
-> 2675         fetched = self._callable_fn(*array_vals)
2676     return fetched[:len(self.outputs)]
2677

```

```

/anaconda3/lib/python3.6/site-packages/tensorflow/python/client/session.py in __call__
1437     ret = tf_session.TF_SessionRunCallable(
1438         self._session._session, self._handle, args, status,
-> 1439         run_metadata_ptr)
1440     if run_metadata:
1441         proto_data = tf_session.TF_GetBuffer(run_metadata_ptr)

```

KeyboardInterrupt:

```

In [14]: model.load_weights('cropping-1.h5')
        loss1 = model.evaluate(val_a, val_b, verbose=0)
        model.load_weights('cropping-2.h5')
        loss2 = model.evaluate(val_a, val_b, verbose=0)
        model.load_weights('cropping-3.h5')
        loss3 = model.evaluate(val_a, val_b, verbose=0)
        model_name = 'cropping-1.h5'
        if loss2 <= loss1 and loss2 < loss3: model_name = 'cropping-2.h5'
        if loss3 <= loss1 and loss3 <= loss2: model_name = 'cropping-3.h5'
        model.load_weights(model_name)
        loss1, loss2, loss3, model_name

```

Out[14]: (26.919307861328125, 'cropping-1.h5')

```

In [15]: model2 = build_model(with_dropout=False)
        model2.load_weights(model_name)
        model2.summary()

```

Layer (type)	Output Shape	Param #	Connected to
input_3 (InputLayer)	(None, 128, 128, 1)	0	
conv2d_35 (Conv2D)	(None, 128, 128, 64)	5248	input_3[0][0]
conv2d_36 (Conv2D)	(None, 128, 128, 64)	36928	conv2d_35[0][0]
batch_normalization_13 (Batch Normalization)	(None, 128, 128, 64)	256	conv2d_36[0][0]
conv2d_37 (Conv2D)	(None, 64, 64, 64)	16448	batch_normalization_13[0][0]

conv2d_38 (Conv2D)	(None, 64, 64, 64)	36928	conv2d_37[0][0]
conv2d_39 (Conv2D)	(None, 64, 64, 64)	36928	conv2d_38[0][0]
batch_normalization_14 (BatchNo	(None, 64, 64, 64)	256	conv2d_39[0][0]
conv2d_40 (Conv2D)	(None, 32, 32, 64)	16448	batch_normalization_14[0][0]
conv2d_41 (Conv2D)	(None, 32, 32, 64)	36928	conv2d_40[0][0]
conv2d_42 (Conv2D)	(None, 32, 32, 64)	36928	conv2d_41[0][0]
batch_normalization_15 (BatchNo	(None, 32, 32, 64)	256	conv2d_42[0][0]
conv2d_43 (Conv2D)	(None, 16, 16, 64)	16448	batch_normalization_15[0][0]
conv2d_44 (Conv2D)	(None, 16, 16, 64)	36928	conv2d_43[0][0]
conv2d_45 (Conv2D)	(None, 16, 16, 64)	36928	conv2d_44[0][0]
batch_normalization_16 (BatchNo	(None, 16, 16, 64)	256	conv2d_45[0][0]
conv2d_46 (Conv2D)	(None, 8, 8, 64)	16448	batch_normalization_16[0][0]
conv2d_47 (Conv2D)	(None, 8, 8, 64)	36928	conv2d_46[0][0]
conv2d_48 (Conv2D)	(None, 8, 8, 64)	36928	conv2d_47[0][0]
batch_normalization_17 (BatchNo	(None, 8, 8, 64)	256	conv2d_48[0][0]
conv2d_49 (Conv2D)	(None, 4, 4, 64)	16448	batch_normalization_17[0][0]
conv2d_50 (Conv2D)	(None, 4, 4, 64)	36928	conv2d_49[0][0]
conv2d_51 (Conv2D)	(None, 4, 4, 64)	36928	conv2d_50[0][0]
batch_normalization_18 (BatchNo	(None, 4, 4, 64)	256	conv2d_51[0][0]
max_pooling2d_5 (MaxPooling2D)	(None, 4, 1, 64)	0	batch_normalization_18[0][0]
max_pooling2d_6 (MaxPooling2D)	(None, 1, 4, 64)	0	batch_normalization_18[0][0]
flatten_5 (Flatten)	(None, 256)	0	max_pooling2d_5[0][0]
flatten_6 (Flatten)	(None, 256)	0	max_pooling2d_6[0][0]
dense_7 (Dense)	(None, 16)	4112	flatten_5[0][0]

dense_8 (Dense)	(None, 16)	4112	flatten_6[0][0]
concatenate_3 (Concatenate)	(None, 32)	0	dense_7[0][0] dense_8[0][0]
dense_9 (Dense)	(None, 4)	132	concatenate_3[0][0]

```

=====
Total params: 503,588
Trainable params: 502,820
Non-trainable params: 768
=====

```

```
In [16]: model2.compile(Adam(lr=0.002), loss='mean_squared_error')
        model2.evaluate(val_a, val_b, verbose=0)
```

```
Out[16]: 26.919307861328125
```

```
In [17]: # Recompute the mean and variance running average without dropout
        for layer in model2.layers:
            if not isinstance(layer, BatchNormalization):
                layer.trainable = False
        model2.compile(Adam(lr=0.002), loss='mean_squared_error')
        model2.fit_generator(TrainingData(), epochs=1, max_queue_size=12, workers=6, verbose=0)
        for layer in model2.layers:
            if not isinstance(layer, BatchNormalization):
                layer.trainable = True
        model2.compile(Adam(lr=0.002), loss='mean_squared_error')
        model2.save('cropping.model')
```

```
Epoch 1/1
500/500 [=====] - 1376s 3s/step - loss: 10.1915 - val_loss: 24.3823
```

```
In [18]: model2.evaluate(val_a, val_b, verbose=0)
```

```
Out[18]: 24.38229362487793
```

```
In [19]: images = []
        for i, (p, coords) in enumerate(val[:25]):
            a = val_a[i:i+1]
            rect1 = val_b[i]
            rect2 = model2.predict(a).squeeze()
            img = array_to_img(a[0]).convert('RGB')
            draw = Draw(img)
            draw.rectangle(rect1, outline='red')
            draw.rectangle(rect2, outline='yellow')
            images.append(img)
        show_whale(images)
```

